High automation, low error rate

7.7%

User prompts out of 454k total events - the rest is autonomous tool activity

1.15%

API error rate across 118k requests - stable platform behavior

11,335 / 20,656

Median vs. mean session tokens - p99 hits 145k, skewed by long outlier sessions

$60 / user

Average cost per engineer over ~60 days ($6k total across 100 users)

100 engineers

Mostly hands-off per session - Claude Code runs autonomously and reliably

Opus earns its cost; Haiku absorbs the volume

Haiku handles the bulk cheaply. 45,979 requests - 39% of all traffic - at just $185 (3.1% of total cost). Its output-to-input ratio is ~0.12x: it scans, reads, and routes. It ingests 76% of all input tokens while generating almost nothing back.

Opus is where money goes. The two Opus models account for ~49,652 requests but $4,258 - 71% of total spend. Output-to-input ratios of 4-5x reveal what that cost buys: generative, multi-step work where the model writes far more than it reads.

Sonnet sits in the middle. 19,698 requests at $1,366 - moderately complex tasks that don't need Opus-level generation but go beyond Haiku's scanning role.

The pattern is deliberate tiering. Cost concentration in Opus reflects output-heavy generation tasks. Haiku cheaply processes the read-heavy majority. Each model earns its place - Haiku on volume, Opus on value per token out.

Total spend: ~$5,809 across 117,000+ API requests.

Bash is the bottleneck

Bash's 7% failure rate and 5.4s average latency make it the single highest reliability and performance risk across all tools.

Bash ⚠ Risk

42,332 runs
Success rate: 93.0% — lowest of all tools
Avg latency: 5,363ms
Max latency: 21,564ms

Read ✓ Fast

45,077 runs
Success rate: 98.5%
Avg latency: 35ms
153x faster than Bash

Edit ✓ Reliable

18,759 runs
Success rate: 98.9%+
Avg latency: sub-2,000ms
Consistent performance

Grep ✓ Reliable

11,342 runs
Success rate: 98.9%+
Avg latency: sub-2,000ms
Consistent performance

Spend scales with headcount, not seniority

<24% spread

tokens per session, L8 to L10 - depth is flat across levels

5x cost difference

L6 ($1,321) vs L10 ($278) - tracks user count directly

100 of 100 users

active at peak hours (11-13, 16-17) - no idle tail

CV >1.38 everywhere

session variability is platform-wide, not practice-specific